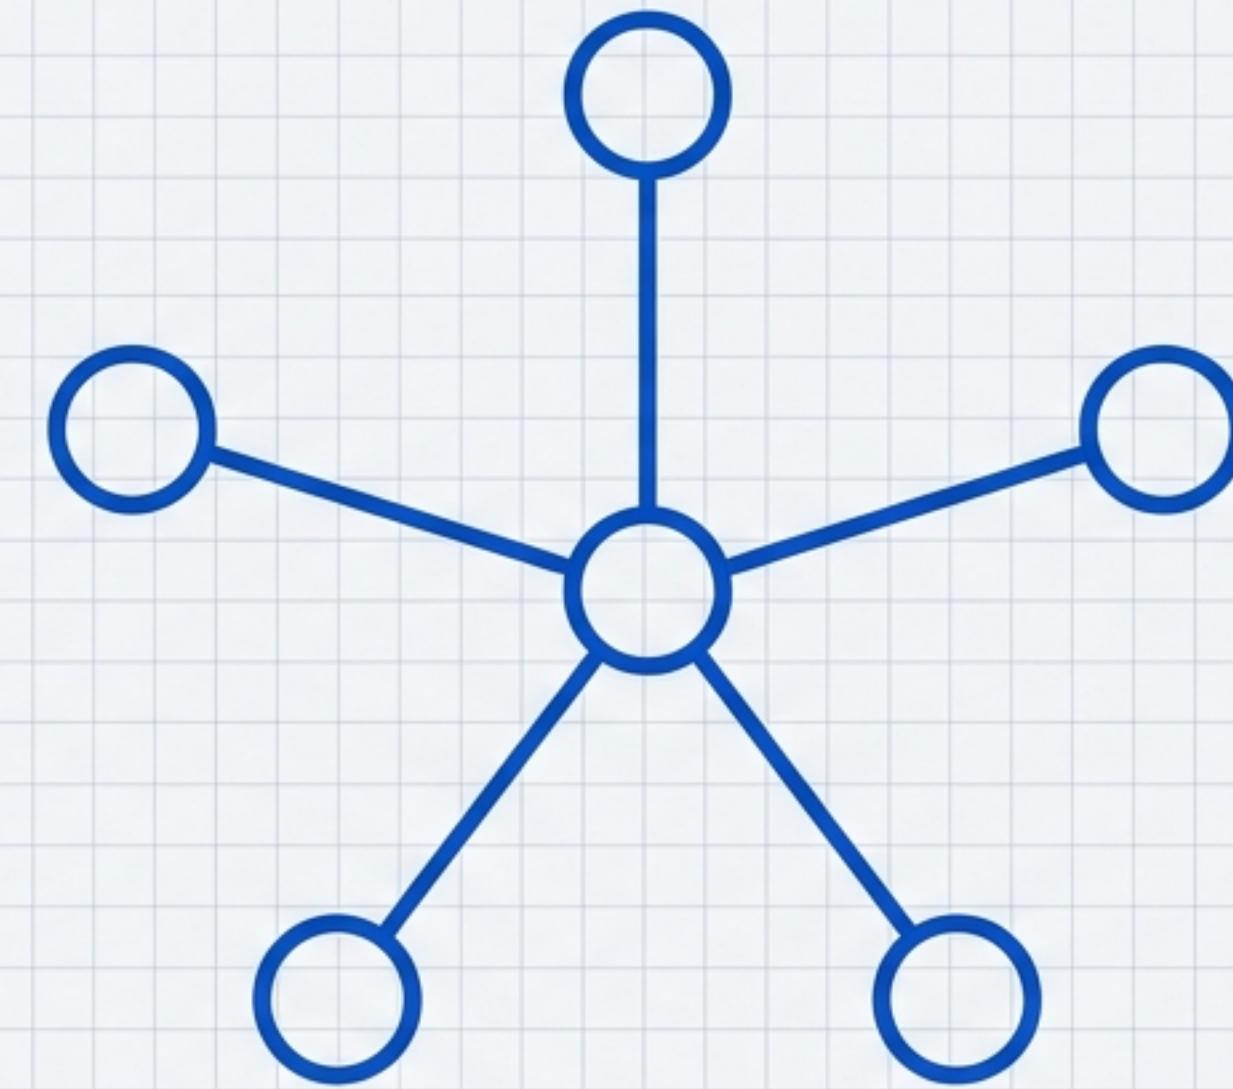


DISTRIBUTED SYSTEMS

Principles & Paradigms:
The Architect's Blueprint

Based on the text by Maarten van
Steen & Andrew S. Tanenbaum



A collection of independent
computers that appears to its users as
a single coherent system.

Distributed Systems (Expansive View)

Concept: Spreading is Sufficient.

Driver: Expansion & Dependability
(e.g., Cloud CDNs, Google Mail).

Goal: Coherence. To the user, it appears as one computer.

Decentralized Systems (Integrative View)

Concept: Spreading is Necessary.

Driver: Administrative Boundaries & Trust (e.g., Blockchain, Federated Learning).

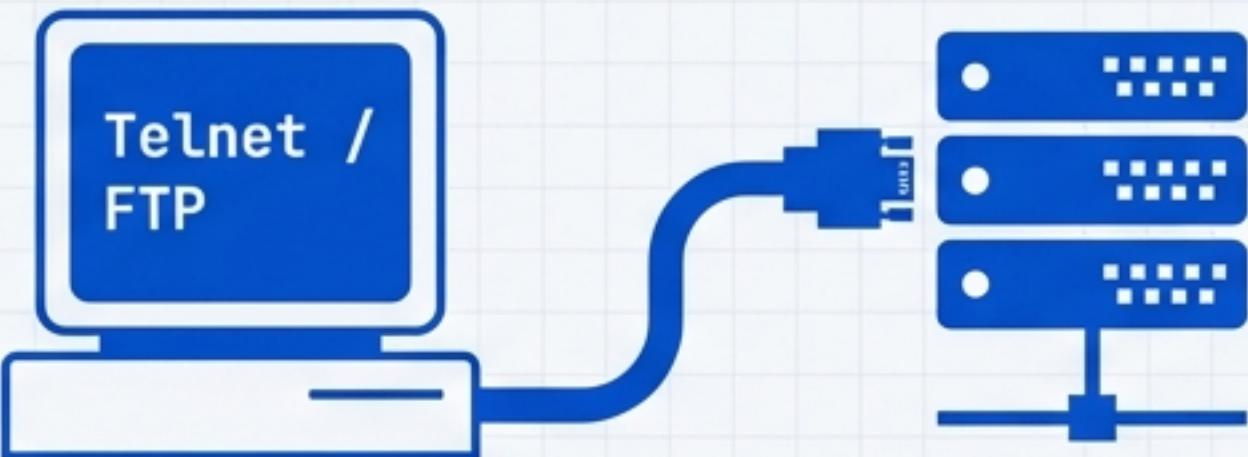
Goal: Integration. The separation is the defining feature.

Key Insight: A distributed system hides its geography; a decentralized system is defined by it.

Design Goal 1: Resource Sharing

From Explicit Connection to Seamless Integration

The Past: Remote Access



Explicit connection to foreign machines.

The Present: Seamless Integration



Cloud Storage & P2P. User places a file locally; system handles the distribution.

Key Drivers

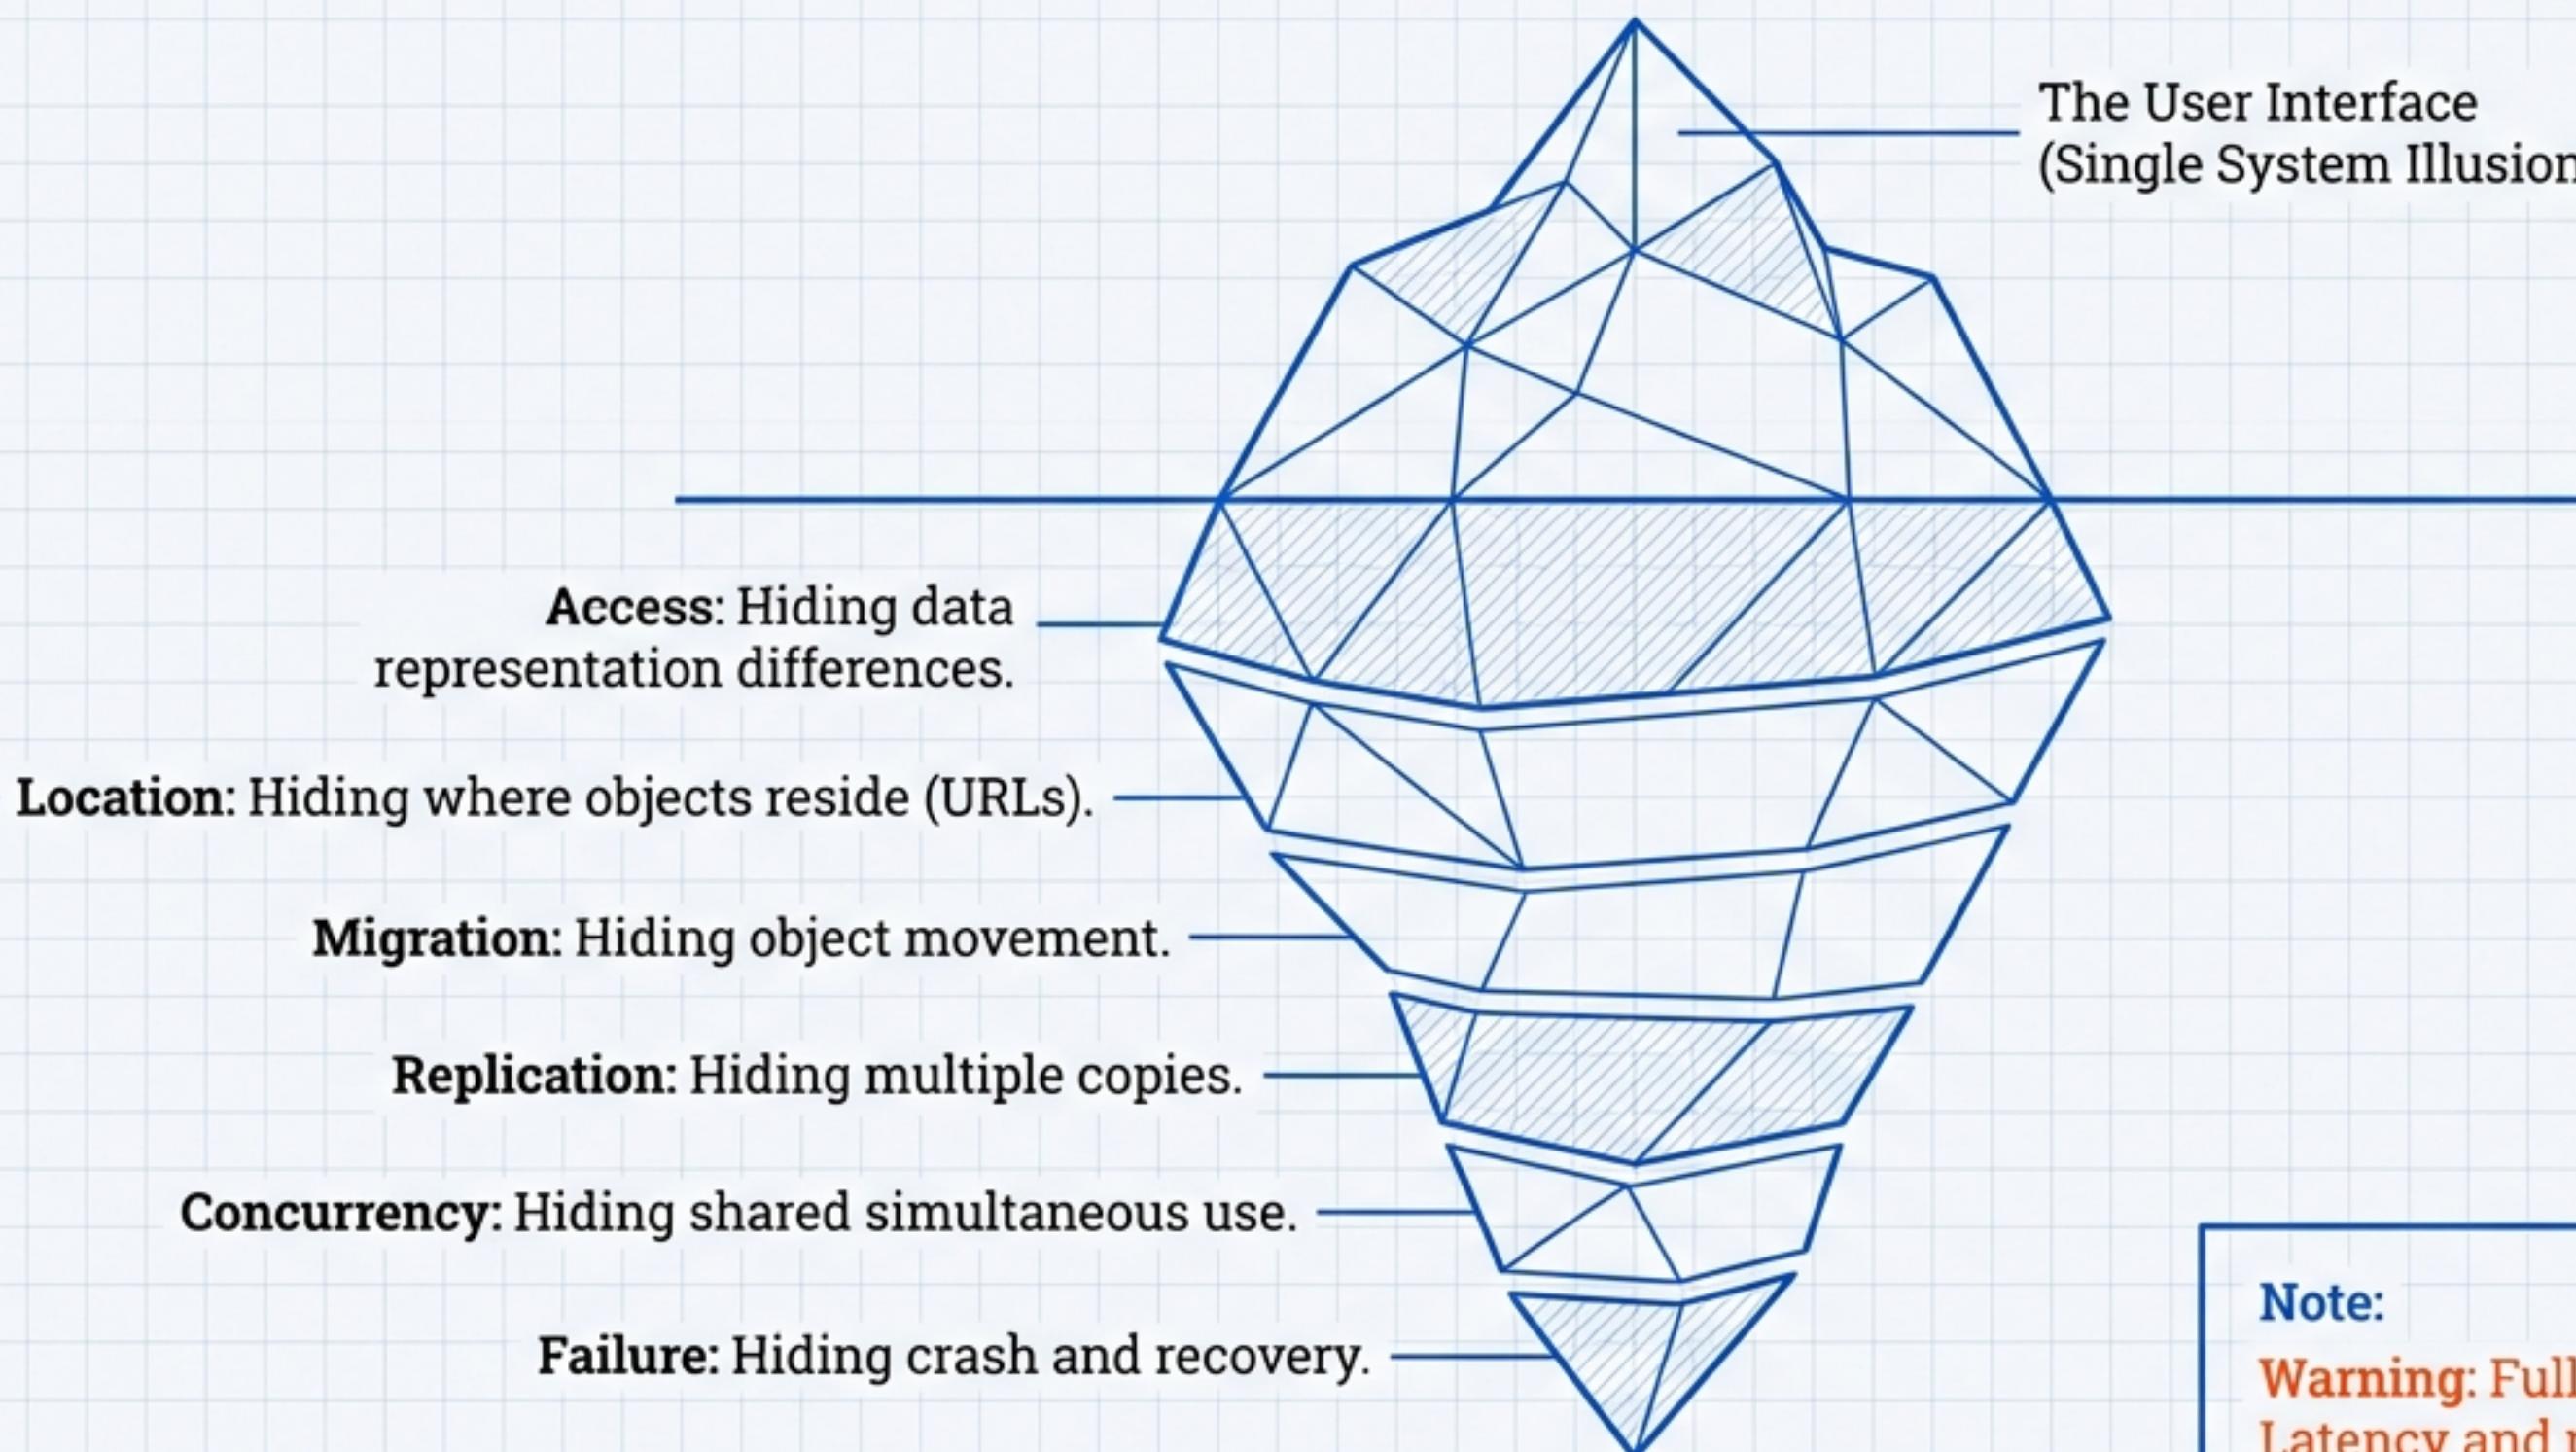
Economics

Sharing expensive peripherals (NAS) and storage infrastructure.

Collaboration

Groupware, collaborative editing, and data exchange.

Design Goal 2: Distribution Transparency



Note:

Warning: Full transparency is impossible.
Latency and physics impose limits.

Design Goal 3: Openness

Separating Policy from Mechanism

Key Requirements



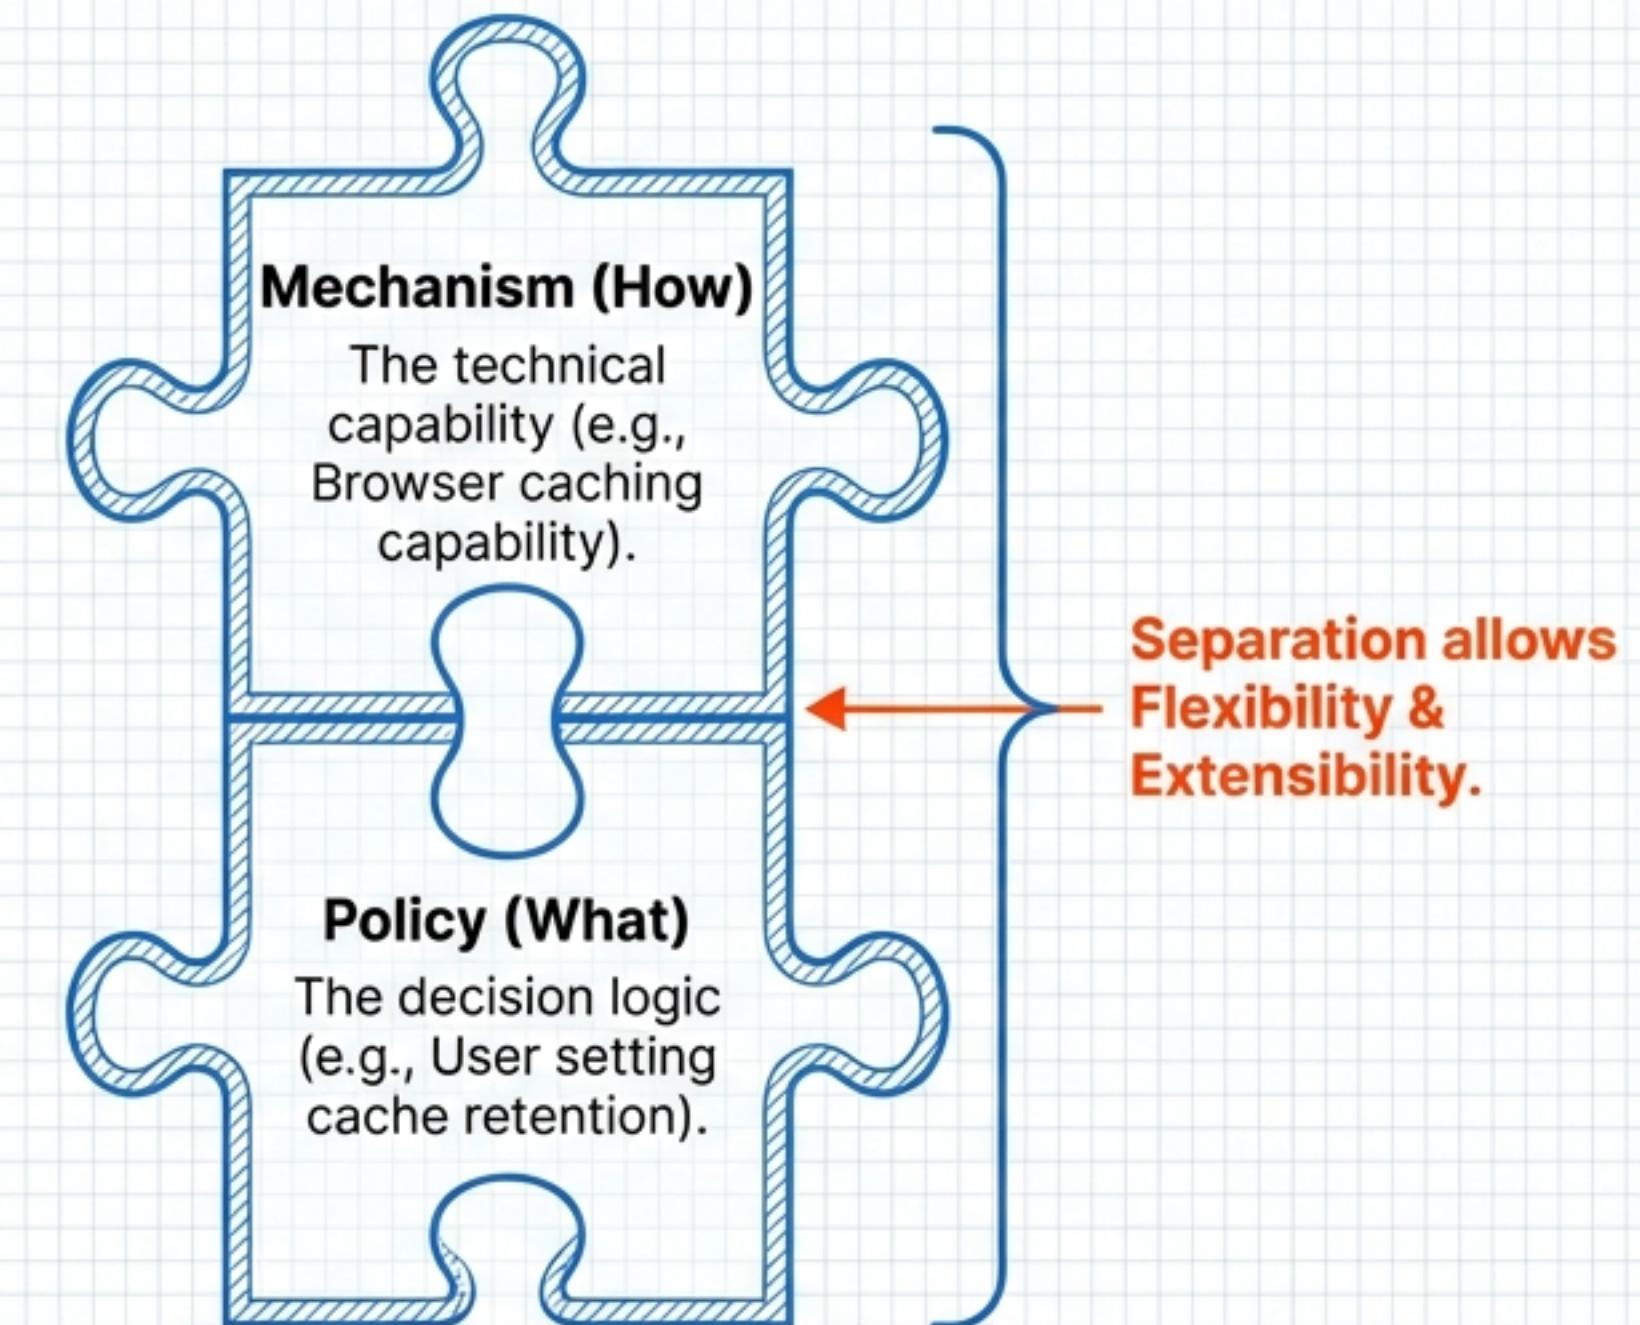
Interoperability: Implementations from different vendors co-existing.



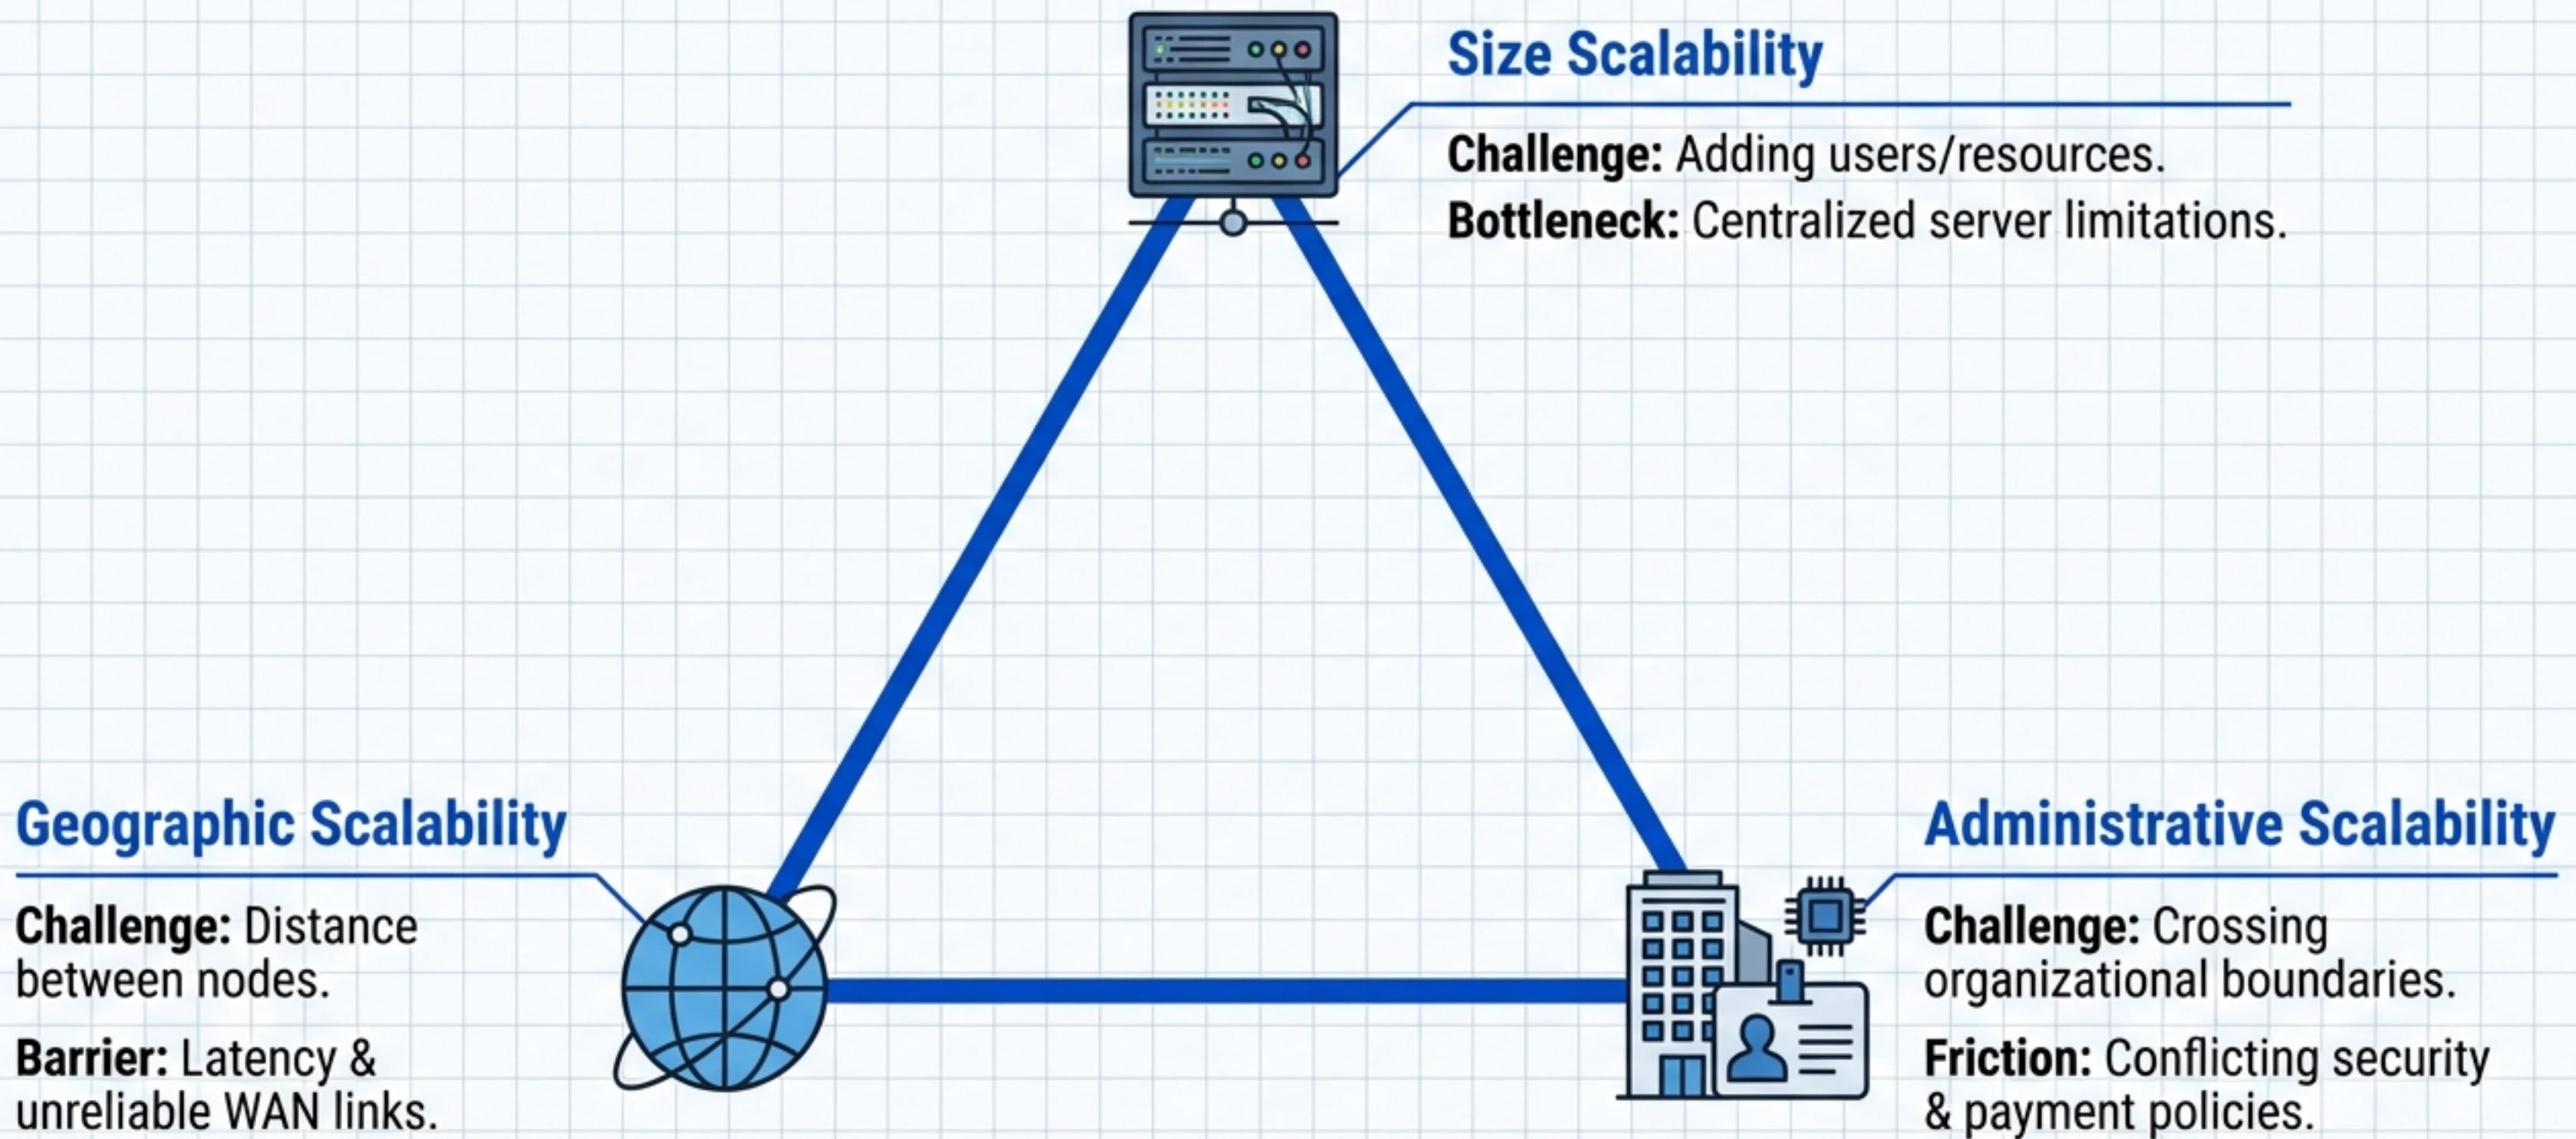
Portability: Application mobility across systems.



IDL (Interface Definition Language): Neutral implementation syntax.



Design Goal 4: The Dimensions of Scalability

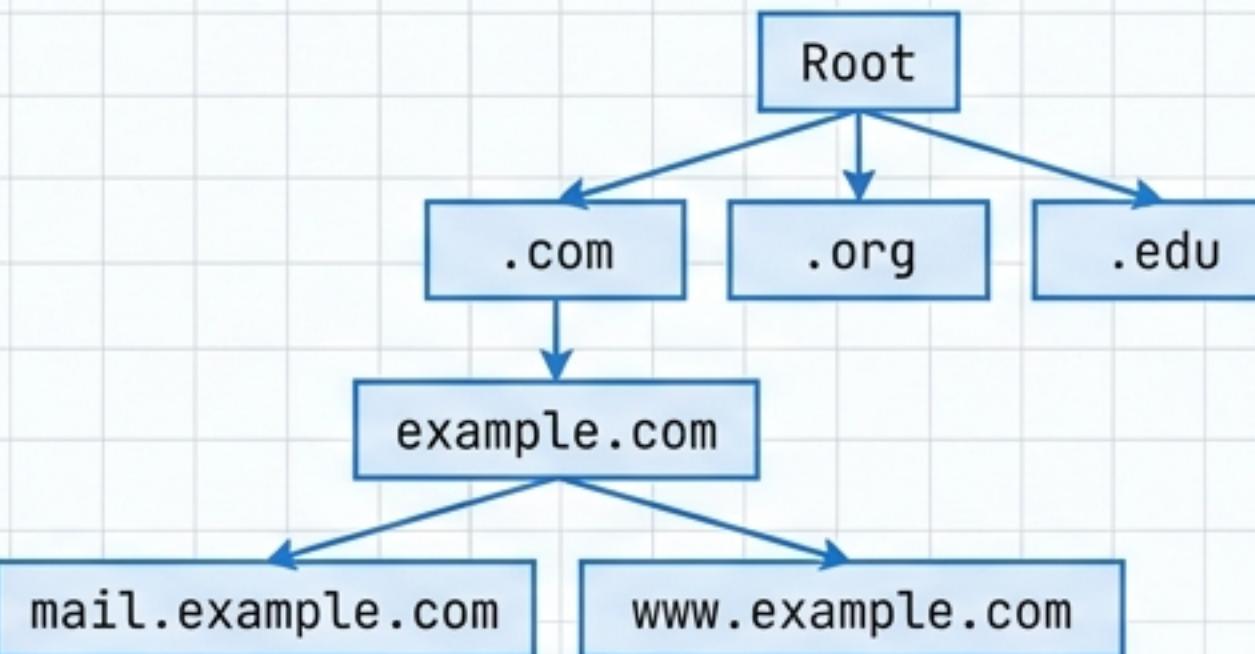


Solving Scalability: Techniques for Expansion

Hiding Latency

- Async Communication: Fire-and-forget instead of blocking.
 - Code Shipping: Move validation logic to the client (e.g., JavaScript).
-

Partitioning



Splitting components into smaller parts (e.g., DNS Zones, Web Sharding).

Replication

- Balancing load and increasing availability.
- **The Trade-off:** Consistency vs. Synchronization costs.

Design Goal 5: Dependability & Fault Tolerance

The Four Pillars of Dependability

Availability:

Ready for immediate use?

Reliability:

Continuous operation without failure?

Safety:

Catastrophe avoidance upon failure.

Maintainability:

Ease of repair.

The Chain of Causality

Fault (The Cause)

Bug, loose wire, cosmic ray.



Error (The State)

Incorrect bit value or internal state.



Failure (The Behavior)

System fails to meet promises.

Goal: Failure Masking. Providing service despite faults.

Design Goal 6: Security

The CIA Triad



Confidentiality

Authorized disclosure only.



Integrity

Authorized alterations only.



Authentication

Identity verification.

Cryptographic Mechanics

Symmetric



Single shared key for encryption/decryption.

Asymmetric (Public-Key)

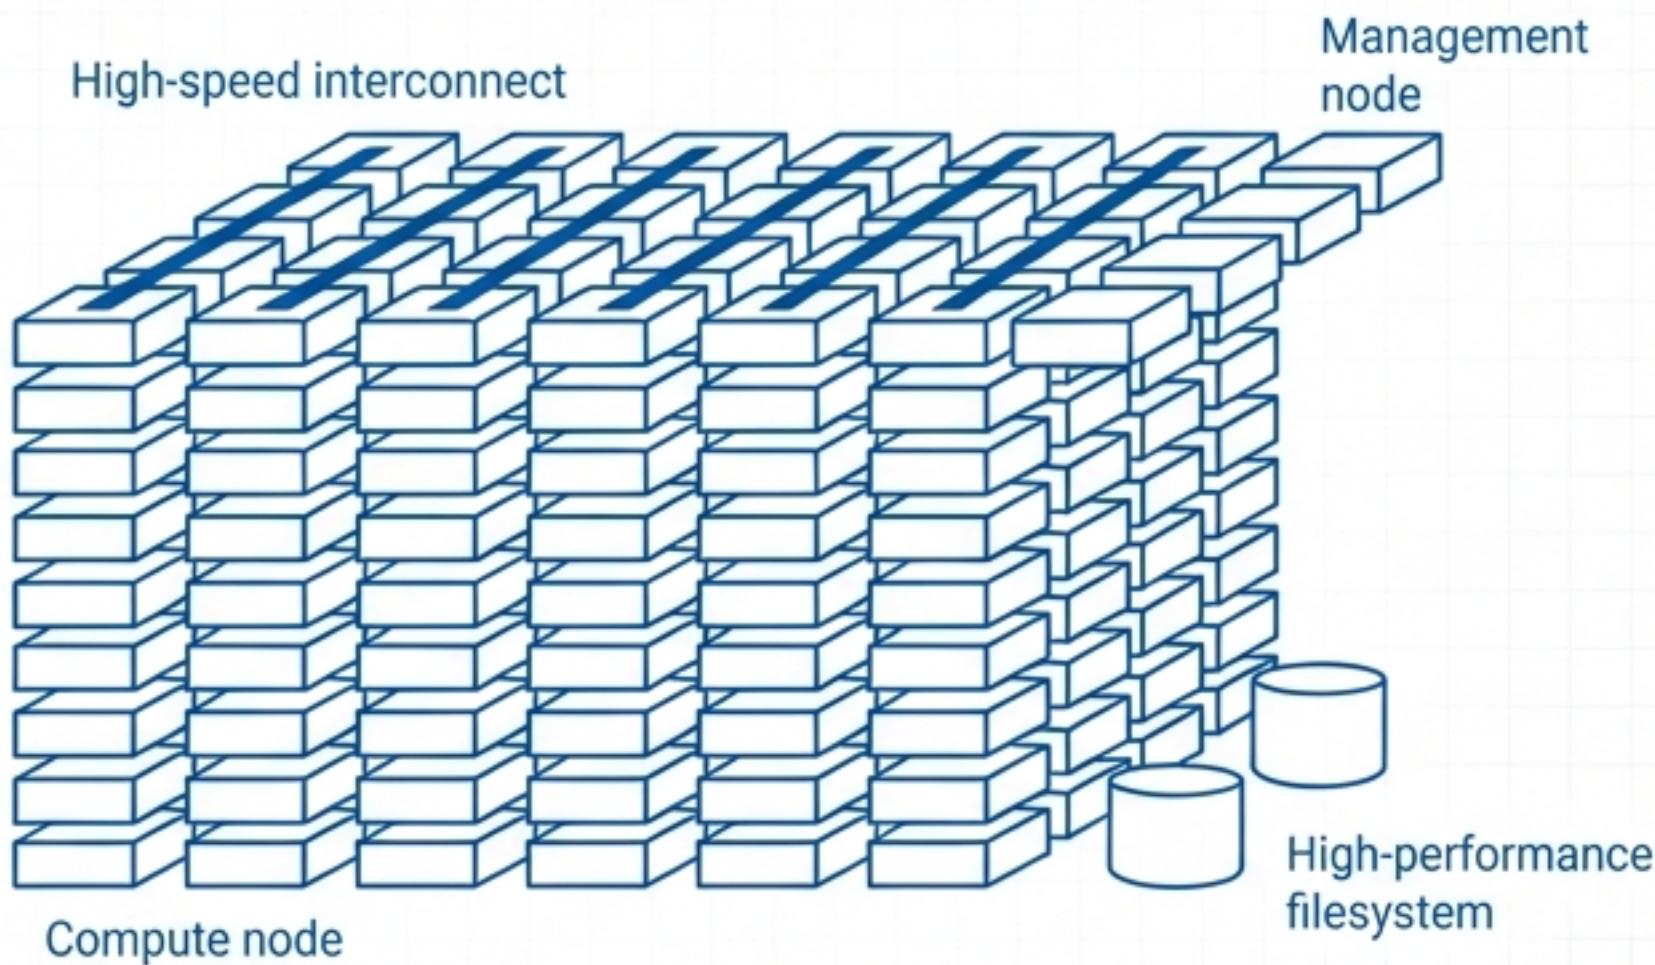


Public key verifies; Private key signs/decrypts

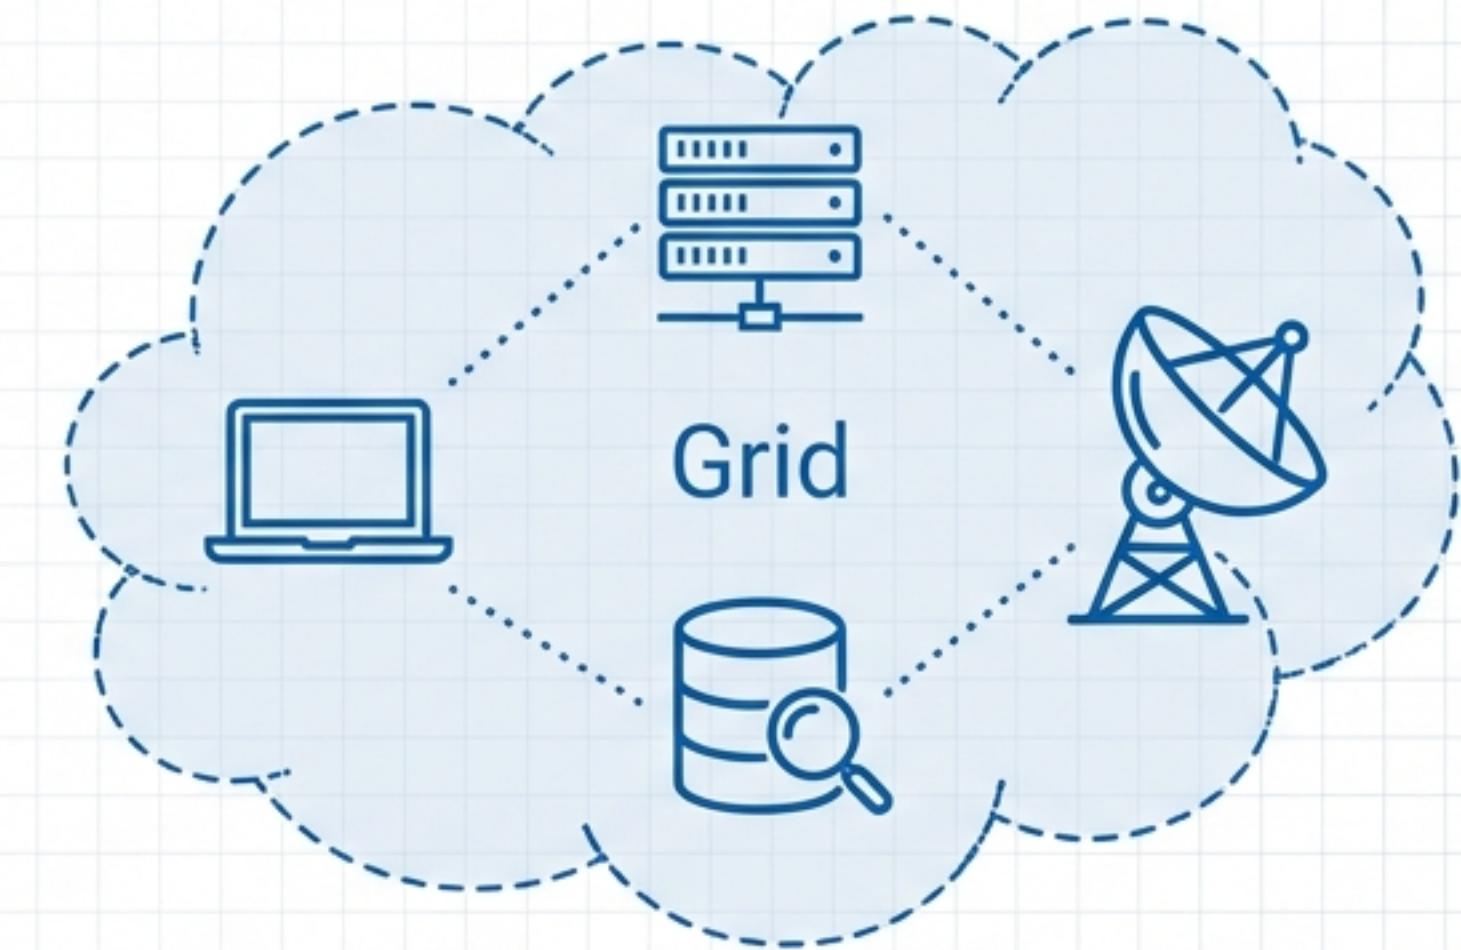
Hashing: Fixed-length strings for integrity verification.

Landscape: High-Performance Computing (HPC)

Cluster Computing



Grid Computing

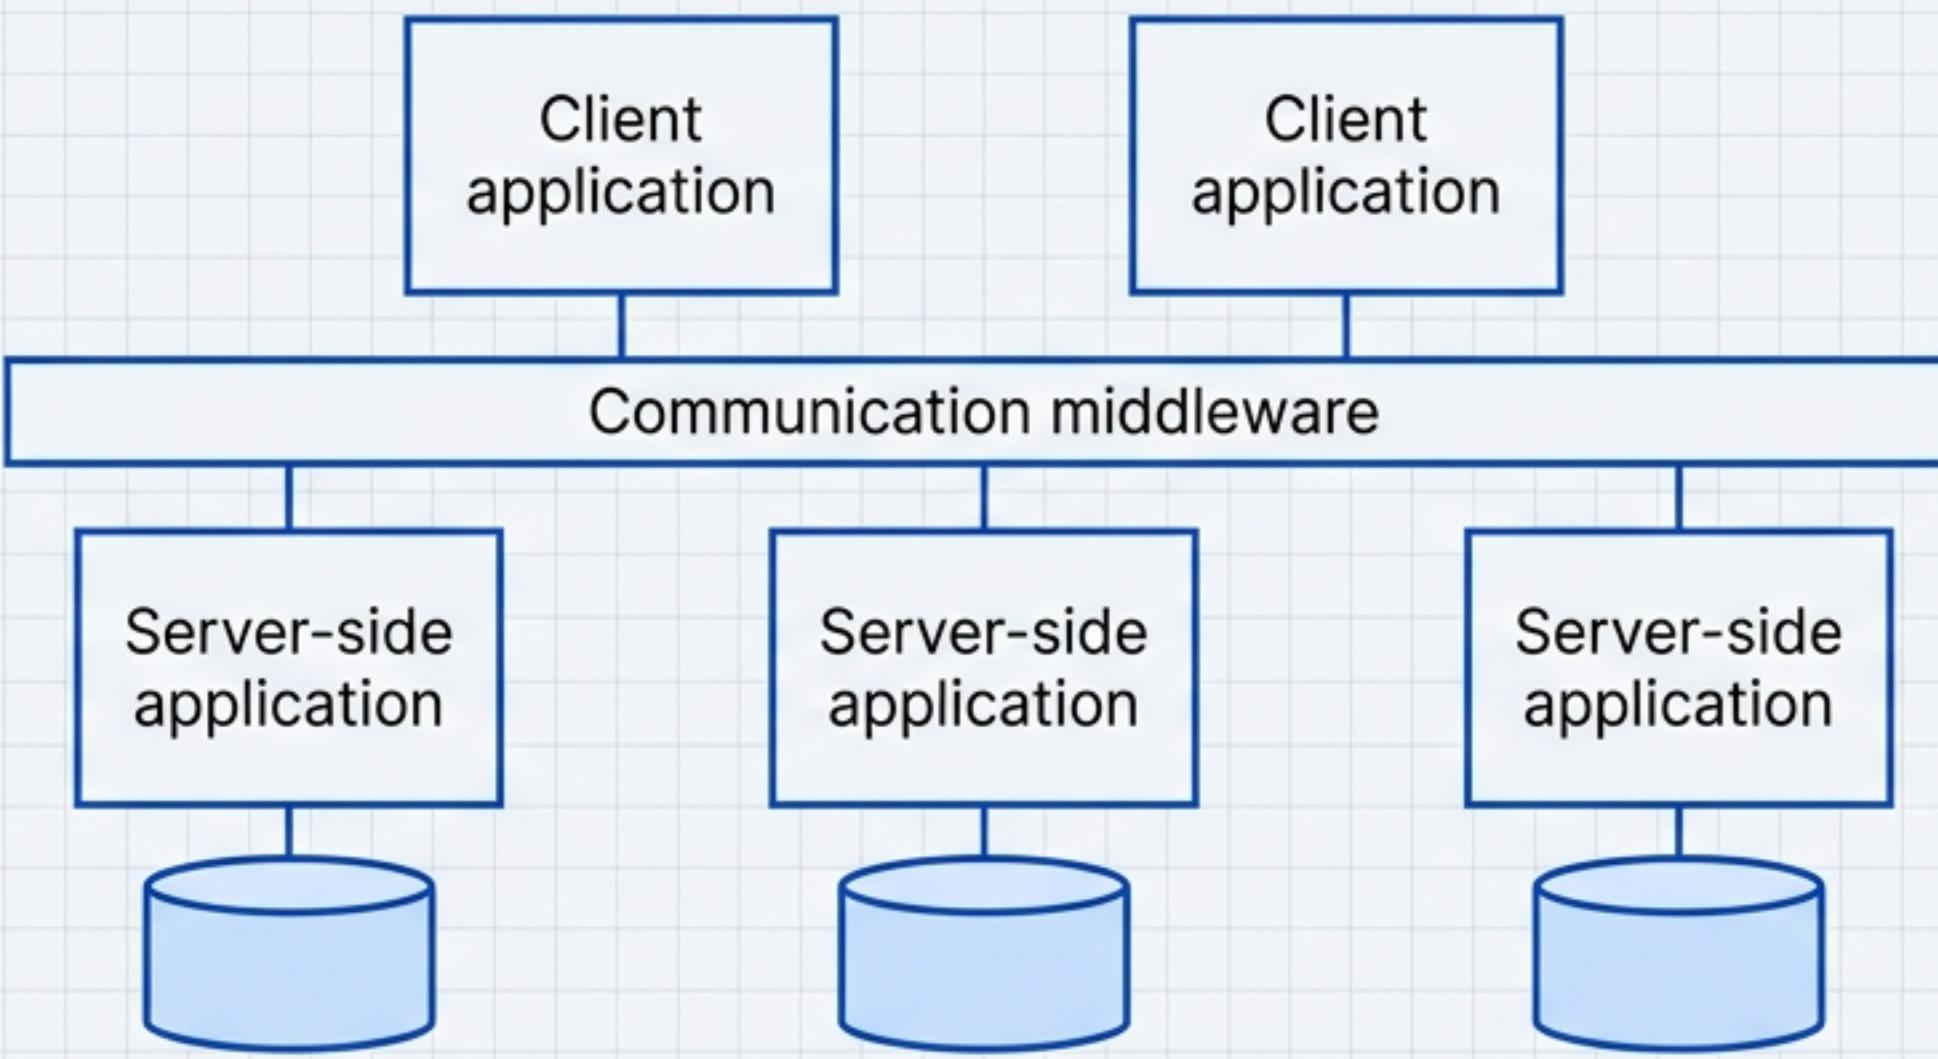


- Homogeneous hardware & OS.
- Single management node.
- High-speed LAN (Tightly coupled).

- Heterogeneous hardware & OS.
- Federated: Spans administrative domains.
- **Virtual Organizations.**

Landscape: Distributed Information Systems

Enterprise Application Integration (EAI)



Transaction Processing (TP)

Operations adhere to **ACID** properties (Atomic, Consistent, Isolated, Durable).

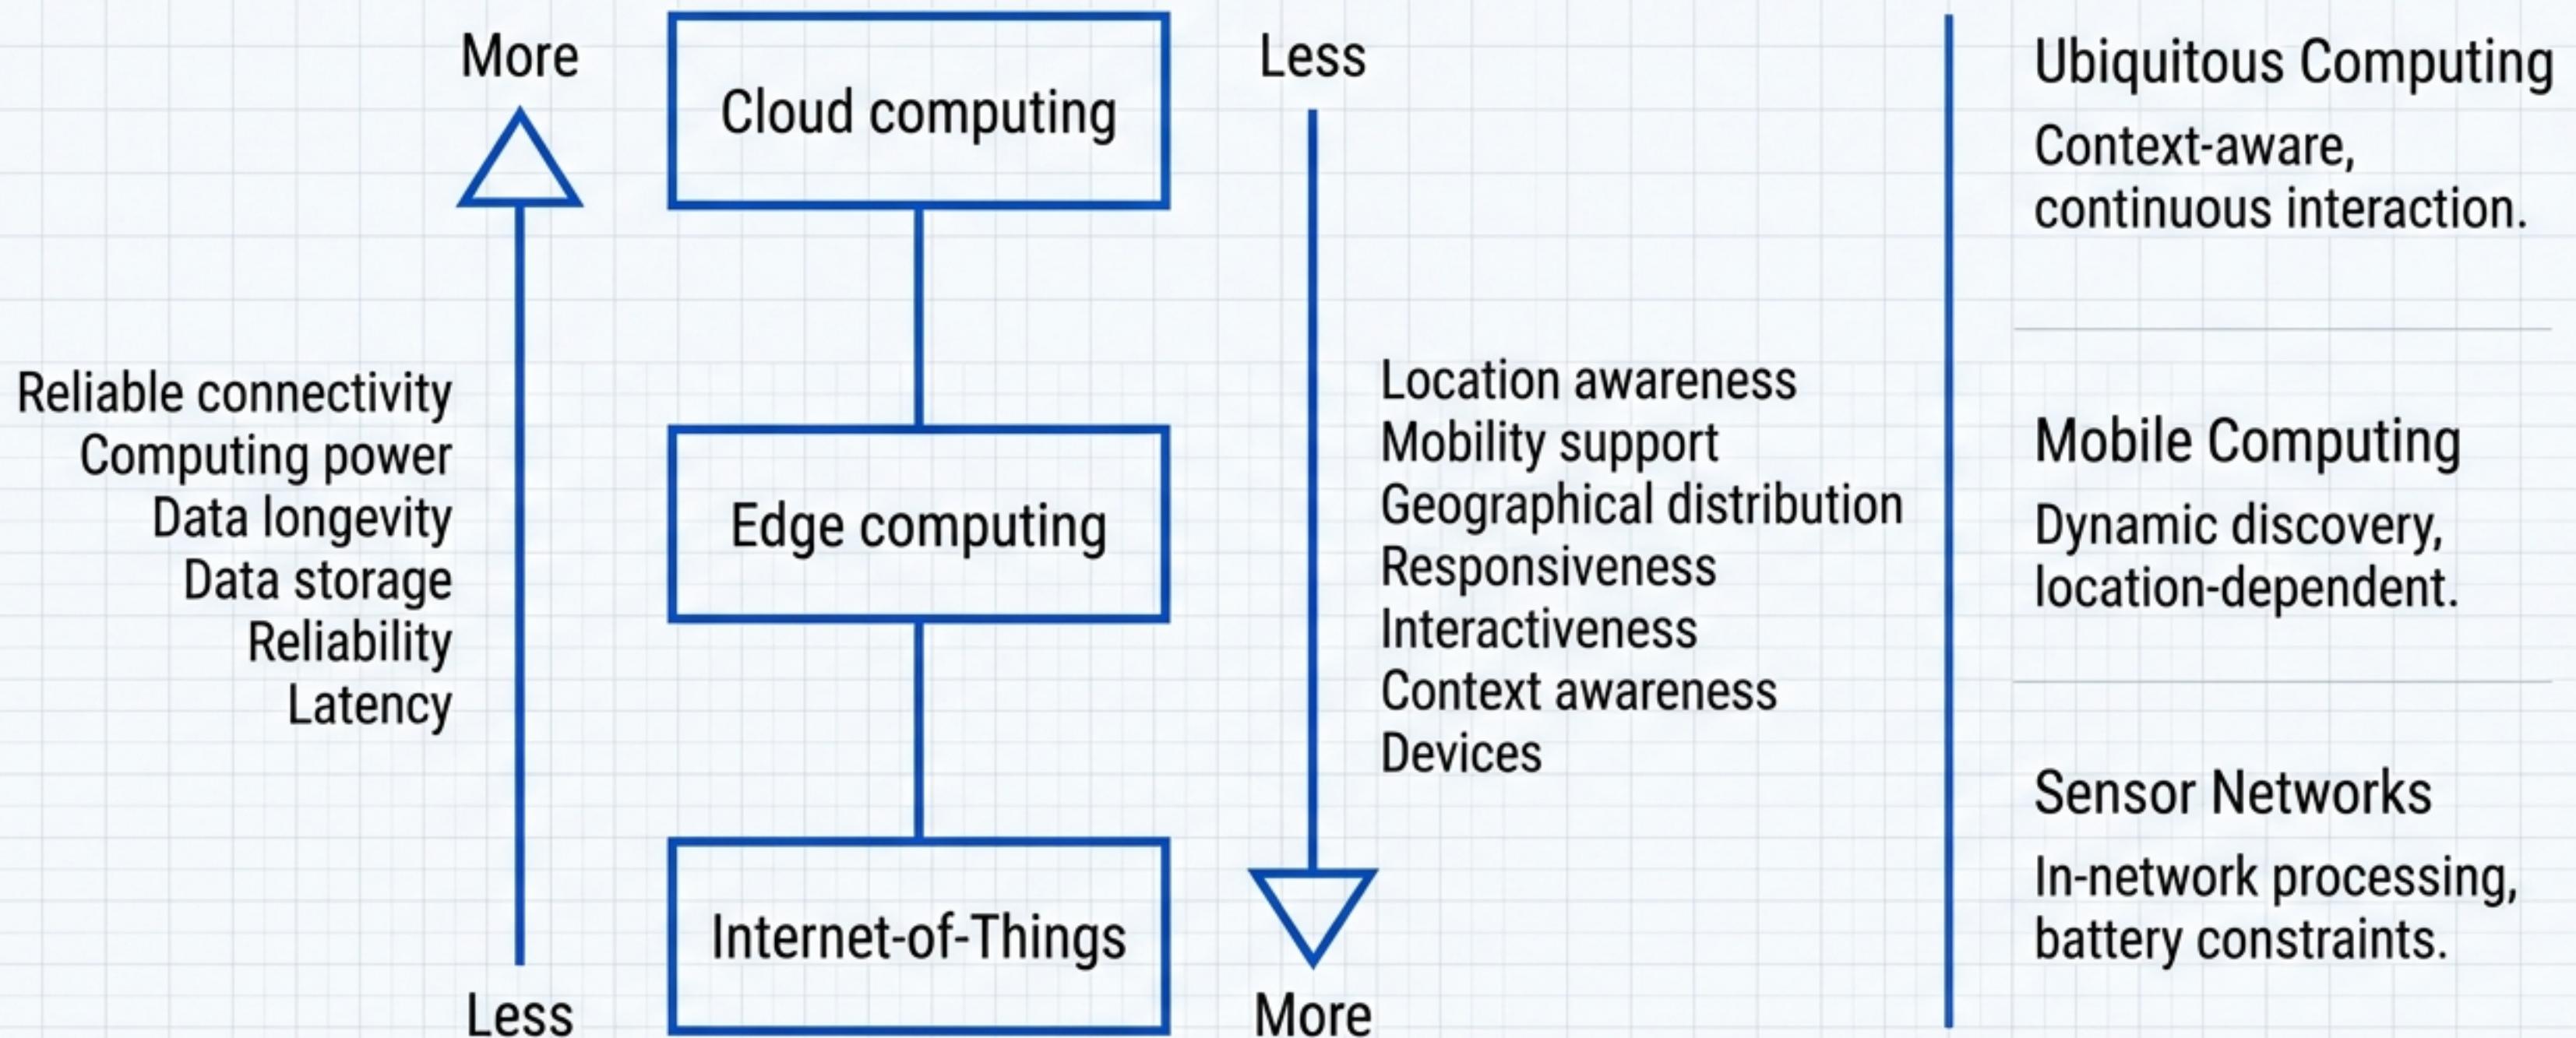
Remote Procedure Call (RPC)

Direct calls between application components.

Message-Oriented Middleware (MOM)

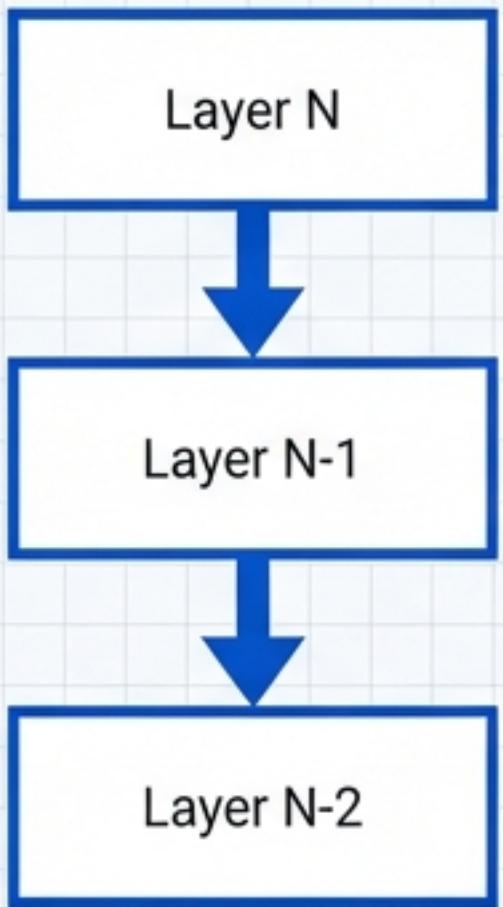
Asynchronous messaging (Publish-Subscribe).

Landscape: Pervasive Systems



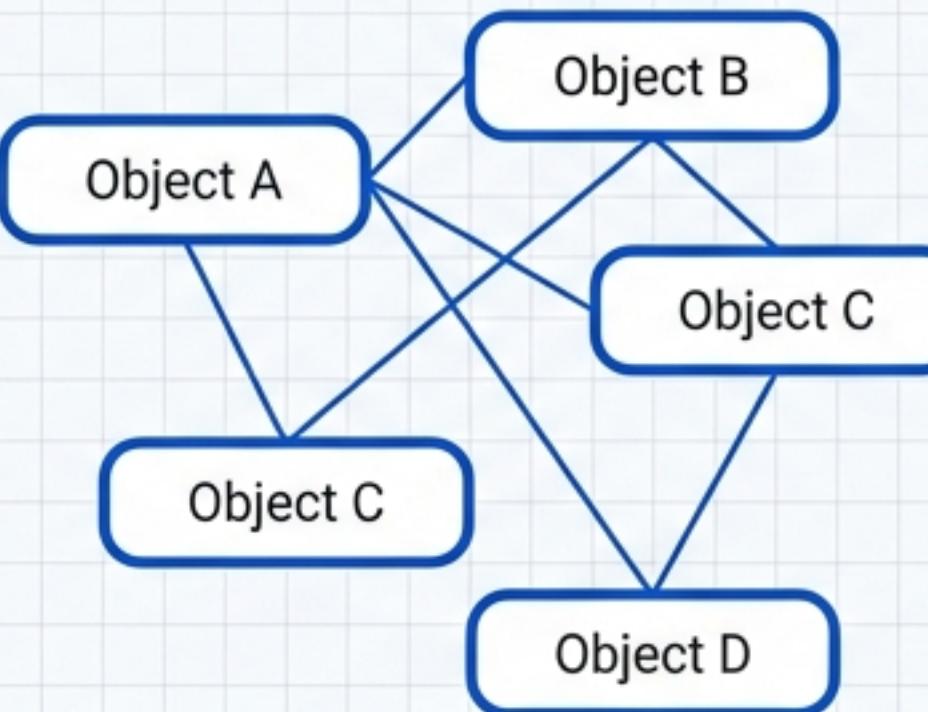
Architectural Styles: Organizing Components

Layered



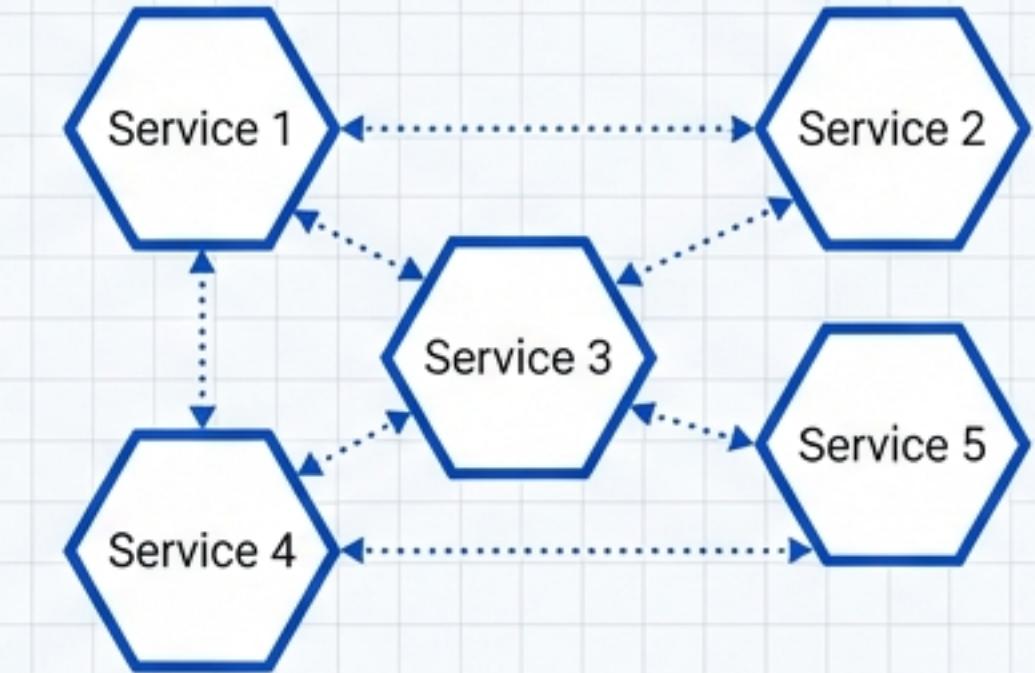
Hierarchical downcalls.
Strict separation.

Object-Based



Roboto Slab
State + Methods encapsulated.
Connected via RMI.

Service-Oriented (SOA)



Roboto Slab
Loosely coupled. "Do one thing well". Independent deployment.

The 8 Pitfalls of Distributed Design

Peter Deutsch's False Assumptions



1	✗ The network is reliable.	2	✗ The network is secure.
3	✗ The network is homogeneous.	4	✗ The topology does not change.
5	✗ Latency is zero.	6	✗ Bandwidth is infinite.
7	✗ Transport cost is zero.	8	✗ There is one administrator.

Successful design requires explicitly managing the fact that all 8 assumptions are false.